

# House Sales in King County, USA

March 22, 2020

```
[1]: import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
import seaborn as sns
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler, PolynomialFeatures
%matplotlib inline
```

## 1.0 Importing the Data

```
[2]: file_name='https://s3-api.us-géo.objectstorage.softlayer.net/cf-courses-data/
→CognitiveClass/DA0101EN/coursera/project/kc_house_data_NaN.csv'
df=pd.read_csv(file_name)
```

```
[3]: df.head()
```

```
[3]: Unnamed: 0      id      date      price  bedrooms  bathrooms  \
0          0  7129300520  20141013T000000  221900.0         3.0         1.00
1          1  6414100192  20141209T000000  538000.0         3.0         2.25
2          2  5631500400  20150225T000000  180000.0         2.0         1.00
3          3  2487200875  20141209T000000  604000.0         4.0         3.00
4          4  1954400510  20150218T000000  510000.0         3.0         2.00
```

```
sqft_living  sqft_lot  floors  waterfront  ...  grade  sqft_above  \
0          1180      5650      1.0           0  ...      7          1180
1          2570      7242      2.0           0  ...      7          2170
2           770     10000      1.0           0  ...      6           770
3          1960      5000      1.0           0  ...      7          1050
4          1680      8080      1.0           0  ...      8          1680
```

```
sqft_basement  yr_built  yr_renovated  zipcode      lat      long  \
0              0      1955            0    98178  47.5112 -122.257
1            400      1951          1991    98125  47.7210 -122.319
2              0      1933            0    98028  47.7379 -122.233
3            910      1965            0    98136  47.5208 -122.393
4              0      1987            0    98074  47.6168 -122.045
```

```
sqft_living15  sqft_lot15
0            1340        5650
```

1	1690	7639
2	2720	8062
3	1360	5000
4	1800	7503

[5 rows x 22 columns]

[4]: df.dtypes

```
[4]: Unnamed: 0      int64
id              int64
date            object
price           float64
bedrooms        float64
bathrooms       float64
sqft_living     int64
sqft_lot        int64
floors          float64
waterfront      int64
view            int64
condition       int64
grade           int64
sqft_above     int64
sqft_basement  int64
yr_built        int64
yr_renovated    int64
zipcode        int64
lat            float64
long           float64
sqft_living15  int64
sqft_lot15     int64
dtype: object
```

[5]: df.describe()

```
[5]:
```

	Unnamed: 0	id	price	bedrooms	bathrooms	\
count	21613.00000	2.161300e+04	2.161300e+04	21600.000000	21603.000000	
mean	10806.00000	4.580302e+09	5.400881e+05	3.372870	2.115736	
std	6239.28002	2.876566e+09	3.671272e+05	0.926657	0.768996	
min	0.00000	1.000102e+06	7.500000e+04	1.000000	0.500000	
25%	5403.00000	2.123049e+09	3.219500e+05	3.000000	1.750000	
50%	10806.00000	3.904930e+09	4.500000e+05	3.000000	2.250000	
75%	16209.00000	7.308900e+09	6.450000e+05	4.000000	2.500000	
max	21612.00000	9.900000e+09	7.700000e+06	33.000000	8.000000	

	sqft_living	sqft_lot	floors	waterfront	view	\
count	21613.000000	2.161300e+04	21613.000000	21613.000000	21613.000000	
mean	2079.899736	1.510697e+04	1.494309	0.007542	0.234303	
std	918.440897	4.142051e+04	0.539989	0.086517	0.766318	

min	290.000000	5.200000e+02	1.000000	0.000000	0.000000
25%	1427.000000	5.040000e+03	1.000000	0.000000	0.000000
50%	1910.000000	7.618000e+03	1.500000	0.000000	0.000000
75%	2550.000000	1.068800e+04	2.000000	0.000000	0.000000
max	13540.000000	1.651359e+06	3.500000	1.000000	4.000000

	...	grade	sqft_above	sqft_basement	yr_built	\
count	...	21613.000000	21613.000000	21613.000000	21613.000000	
mean	...	7.656873	1788.390691	291.509045	1971.005136	
std	...	1.175459	828.090978	442.575043	29.373411	
min	...	1.000000	290.000000	0.000000	1900.000000	
25%	...	7.000000	1190.000000	0.000000	1951.000000	
50%	...	7.000000	1560.000000	0.000000	1975.000000	
75%	...	8.000000	2210.000000	560.000000	1997.000000	
max	...	13.000000	9410.000000	4820.000000	2015.000000	

	yr_renovated	zipcode	lat	long	sqft_living15	\
count	21613.000000	21613.000000	21613.000000	21613.000000	21613.000000	
mean	84.402258	98077.939805	47.560053	-122.213896	1986.552492	
std	401.679240	53.505026	0.138564	0.140828	685.391304	
min	0.000000	98001.000000	47.155900	-122.519000	399.000000	
25%	0.000000	98033.000000	47.471000	-122.328000	1490.000000	
50%	0.000000	98065.000000	47.571800	-122.230000	1840.000000	
75%	0.000000	98118.000000	47.678000	-122.125000	2360.000000	
max	2015.000000	98199.000000	47.777600	-121.315000	6210.000000	

	sqft_lot15
count	21613.000000
mean	12768.455652
std	27304.179631
min	651.000000
25%	5100.000000
50%	7620.000000
75%	10083.000000
max	871200.000000

[8 rows x 21 columns]

## 2.0 Data Wrangling

```
[6]: df.drop(['id', 'Unnamed: 0'], axis = 1, inplace = True)
df.describe()
```

	price	bedrooms	bathrooms	sqft_living	sqft_lot	\
count	2.161300e+04	21600.000000	21603.000000	21613.000000	2.161300e+04	
mean	5.400881e+05	3.372870	2.115736	2079.899736	1.510697e+04	
std	3.671272e+05	0.926657	0.768996	918.440897	4.142051e+04	
min	7.500000e+04	1.000000	0.500000	290.000000	5.200000e+02	
25%	3.219500e+05	3.000000	1.750000	1427.000000	5.040000e+03	

50%	4.500000e+05	3.000000	2.250000	1910.000000	7.618000e+03
75%	6.450000e+05	4.000000	2.500000	2550.000000	1.068800e+04
max	7.700000e+06	33.000000	8.000000	13540.000000	1.651359e+06

	floors	waterfront	view	condition	grade \
count	21613.000000	21613.000000	21613.000000	21613.000000	21613.000000
mean	1.494309	0.007542	0.234303	3.409430	7.656873
std	0.539989	0.086517	0.766318	0.650743	1.175459
min	1.000000	0.000000	0.000000	1.000000	1.000000
25%	1.000000	0.000000	0.000000	3.000000	7.000000
50%	1.500000	0.000000	0.000000	3.000000	7.000000
75%	2.000000	0.000000	0.000000	4.000000	8.000000
max	3.500000	1.000000	4.000000	5.000000	13.000000

	sqft_above	sqft_basement	yr_built	yr_renovated	zipcode \
count	21613.000000	21613.000000	21613.000000	21613.000000	21613.000000
mean	1788.390691	291.509045	1971.005136	84.402258	98077.939805
std	828.090978	442.575043	29.373411	401.679240	53.505026
min	290.000000	0.000000	1900.000000	0.000000	98001.000000
25%	1190.000000	0.000000	1951.000000	0.000000	98033.000000
50%	1560.000000	0.000000	1975.000000	0.000000	98065.000000
75%	2210.000000	560.000000	1997.000000	0.000000	98118.000000
max	9410.000000	4820.000000	2015.000000	2015.000000	98199.000000

	lat	long	sqft_living15	sqft_lot15
count	21613.000000	21613.000000	21613.000000	21613.000000
mean	47.560053	-122.213896	1986.552492	12768.455652
std	0.138564	0.140828	685.391304	27304.179631
min	47.155900	-122.519000	399.000000	651.000000
25%	47.471000	-122.328000	1490.000000	5100.000000
50%	47.571800	-122.230000	1840.000000	7620.000000
75%	47.678000	-122.125000	2360.000000	10083.000000
max	47.777600	-121.315000	6210.000000	871200.000000

```
[7]: print("number of NaN values for the column bedrooms :", df['bedrooms'].isnull().
      →sum())
      print("number of NaN values for the column bathrooms :", df['bathrooms'].
      →isnull().sum())
```

number of NaN values for the column bedrooms : 13  
 number of NaN values for the column bathrooms : 10

```
[8]: #replace the missing values of the column 'bedrooms' with the mean of the
      →column 'bedrooms' using the method replace.
      mean=df['bedrooms'].mean()
      df['bedrooms'].replace(np.nan,mean, inplace=True)
```

```
[9]: mean=df['bathrooms'].mean()
df['bathrooms'].replace(np.nan,mean, inplace=True)

[10]: print("number of NaN values for the column bedrooms :", df['bedrooms'].isnull().
        ↳sum())
print("number of NaN values for the column bathrooms :", df['bathrooms'].
        ↳isnull().sum())
```

number of NaN values for the column bedrooms : 0  
number of NaN values for the column bathrooms : 0


### 3.0 Exploratory data analysis

```
[11]: #Use the method value_counts to count the number of houses with unique floor
        ↳values.
floor_values = df['floors'].value_counts()
floor_values.to_frame()
```

```
[11]:      floors
1.0    10680
2.0     8241
1.5     1910
3.0      613
2.5      161
3.5         8
```


```
[12]: #Use the function boxplot to determine whether houses with a waterfront view or
        ↳without have more price outliers.
sns.boxplot(x=df['waterfront'], y = df['price'])
# So houses with a waterfront has more price outliers
```

```
[12]: <matplotlib.axes._subplots.AxesSubplot at 0x1a1e381d30>
```



```
[13]: #Use the function regplot library to determine if the feature sqft_above is
      ↪negatively or positively correlated with price.
      sns.regplot(x=df['sqft_above'], y=df['price'])
      # they are positively correlated with price
```

```
[13]: <matplotlib.axes._subplots.AxesSubplot at 0x1a1f40e438>
```



```
[14]: df.corr()['price'].sort_values()
```

```
[14]: zipcode      -0.053203
      long         0.021626
      condition    0.036362
      yr_built     0.054012
      sqft_lot15    0.082447
      sqft_lot      0.089661
      yr_renovated  0.126434
      floors       0.256794
      waterfront   0.266369
      lat          0.307003
      bedrooms     0.308797
      sqft_basement 0.323816
      view         0.397293
      bathrooms    0.525738
      sqft_living15 0.585379
      sqft_above   0.605567
      grade        0.667434
      sqft_living  0.702035
      price        1.000000
      Name: price, dtype: float64
```

#### Module 4: Model Development

```
[ ]: #Linear Regression
```

```

[15]: import matplotlib.pyplot as plt
      from sklearn.linear_model import LinearRegression

[20]: features = ["floors", "waterfront","lat" ,"bedrooms" ,"sqft_basement" ,"view"
      ↪,"bathrooms","sqft_living15","sqft_above","grade","sqft_living"]

[21]: features = df[features]
      y1 = df['price']

[22]: lm = LinearRegression()
      lm.fit(features, y1)
      lm.score(features, y1)

[22]: 0.657679183672129

[23]: #create a pipeline to build polynomial regression

[24]: Input=[('scale',StandardScaler()),('polynomial',
      ↪PolynomialFeatures(include_bias=False)),('model',LinearRegression())]

[25]: pipe=Pipeline(Input)
      pipe

[25]: Pipeline(memory=None,
      steps=[('scale',
              StandardScaler(copy=True, with_mean=True, with_std=True)),
              ('polynomial',
               PolynomialFeatures(degree=2, include_bias=False,
                                   interaction_only=False, order='C')),
              ('model',
               LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None,
                                 normalize=False))],
      verbose=False)

[26]: pipe.fit(features,y1)

[26]: Pipeline(memory=None,
      steps=[('scale',
              StandardScaler(copy=True, with_mean=True, with_std=True)),
              ('polynomial',
               PolynomialFeatures(degree=2, include_bias=False,
                                   interaction_only=False, order='C')),
              ('model',
               LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None,
                                 normalize=False))],
      verbose=False)

[27]: pipe.score(features, y1)

[27]: 0.7513408553309376

[28]: # We can see polynomial regression has better performance than linear
      ↪regression

```



## Module 5: MODEL EVALUATION AND REFINEMENT

```
[29]: from sklearn.model_selection import cross_val_score
      from sklearn.model_selection import train_test_split
```

```
[32]: #split data to training and testing
      features = ["floors", "waterfront","lat" ,"bedrooms" ,"sqft_basement" ,"view"
      ↪,"bathrooms","sqft_living15","sqft_above","grade","sqft_living"]
      X = df[features ]
      Y = df['price']

      x_train, x_test, y_train, y_test = train_test_split(X, Y, test_size=0.15,
      ↪random_state=1)

      print("number of test samples :", x_test.shape[0])
      print("number of training samples:",x_train.shape[0])
```

```
number of test samples : 3242
number of training samples: 18371
```

```
[33]: #Create and fit a Ridge regression object using the training data,
      #setting the regularization parameter to 0.1 and calculate the  $R^2$  using the
      ↪test data.

      from sklearn.linear_model import Ridge
      Ridgemodel = Ridge(alpha=0.1)
      Ridgemodel.fit(x_train, y_train)
      Ridgemodel.score(x_test, y_test)
```

```
[33]: 0.6478759163939121
```

```
[34]: #Perform a second order polynomial transform
      pr = PolynomialFeatures(degree = 2)
      x_train_pr = pr.fit_transform(x_train)
      x_test_pr = pr.fit_transform(x_test)
      Ridgemodel_pr = Ridge(alpha=0.1)
      Ridgemodel_pr.fit(x_train_pr, y_train)
      Ridgemodel_pr.score(x_test_pr, y_test)
```

```
[34]: 0.7002744279699229
```

```
[35]: #To choose the best alpha
      from sklearn.model_selection import GridSearchCV
```

```
[36]: parameter2 = [{'alpha':[0.01,0.1,1,10,100],'normalize':[True,False]]}
      RR = Ridge()
      Grid1 = GridSearchCV(RR, parameter2, cv=4)
      Grid1.fit(x_train_pr, y_train)
      Grid1.best_estimator_
```

```
//anaconda3/lib/python3.7/site-packages/sklearn/linear_model/ridge.py:147:
LinAlgWarning: Ill-conditioned matrix (rcond=1.85665e-17): result may not be
```

```

accurate.
    overwrite_a=True).T
//anaconda3/lib/python3.7/site-packages/sklearn/linear_model/ridge.py:147:
LinAlgWarning: Ill-conditioned matrix (rcond=2.27951e-17): result may not be
accurate.
    overwrite_a=True).T
//anaconda3/lib/python3.7/site-packages/sklearn/linear_model/ridge.py:147:
LinAlgWarning: Ill-conditioned matrix (rcond=1.99368e-17): result may not be
accurate.
    overwrite_a=True).T
//anaconda3/lib/python3.7/site-packages/sklearn/linear_model/ridge.py:147:
LinAlgWarning: Ill-conditioned matrix (rcond=1.48596e-17): result may not be
accurate.
    overwrite_a=True).T

```

```

[36]: Ridge(alpha=0.01, copy_X=True, fit_intercept=True, max_iter=None,
        normalize=False, random_state=None, solver='auto', tol=0.001)

```

```

[37]: # we use the best alpha '0.01' to build another model
pr = PolynomialFeatures(degree = 2)
x_train_pr = pr.fit_transform(x_train)
x_test_pr = pr.fit_transform(x_test)
Ridgemodel_pr = Ridge(alpha=0.01)
Ridgemodel_pr.fit(x_train_pr, y_train)
Ridgemodel_pr.score(x_test_pr, y_test)

```

```

[37]: 0.7017148826621857

```

```

[:]: # We get the largest R^2 here.
# Here 70.17% data can be explained with this new model.

```